

Last Time:

- Constrained Optimization
- Augmented Lagrangian
- QPs
- Regularization

Today:

- Merit functions + Line Searches
- Control History
- Deterministic Optimal Control

* Merit Functions

- How do we do a line search on a root-finding problem?
- find x^* such that $C(x^*) = 0$
- Define scalar "merit function" $P(x)$ that measures distance to solution
- Standard choices:

$$P(x) = \frac{1}{2} (Cx)^T C(x) = \frac{1}{2} \|Cx\|_2^2$$

$$P(x) = \|Cx\|_1 \quad (\text{any norm works})$$

- Now just do Armijo on $P(x)$:

$$\alpha = 1$$

$$\text{while } P(x + \alpha \bar{x}) > P(x) + b\alpha \nabla P(x)^T \bar{x}$$

$$\alpha \leftarrow \theta \alpha$$

tolerance

step length

end

$$x \leftarrow x + \alpha \Delta x$$

- How about constrained optimization?

$$\begin{array}{ll} \min_x & f(x) \\ \text{s.t.} & g(x) \leq 0 \\ & d(x) = 0 \end{array} \quad \left\{ \begin{array}{l} L(x, \lambda, M) = f(x) + \lambda^T g(x) + M^T d(x) \end{array} \right.$$

- Lots of options for merit functions:

$$P(x, \lambda, M) = \frac{1}{2} \| r_{n+1}(x, \lambda, M) \|_2^2$$

HRT Residual = $\begin{bmatrix} D_x L \\ \min(0, g(x)) \\ d(x) \end{bmatrix}$

$$P(x, \lambda, M) = f(x) + \rho \| \begin{bmatrix} \min(0, g(x)) \\ d(x) \end{bmatrix} \|_1$$

scalar weight

any norm works

$$P(x, \lambda, M) = L_P(x, \lambda, M) \quad (\text{Augmented Lagrangian})$$

* Example Take Aways:

- Excessive constraint penalties can cause problems
- All methods come with a merit function for free

→ Deterministic Optimal Control

- Continuous time:

$$\min_{\begin{cases} x(t) \\ u(t) \end{cases}} J(x(t), u(t)) = \int_{t_0}^{t_f} l(x(t), u(t)) dt + \underline{l_F(x(t_f))}$$

State+input
 trajectories

cost
 function

"Stage Cost"

"terminal cost"

s.t. $\dot{x}(t) = f(x(t), u(t))$ "dynamics constraints"

(possibly other constraints)

- This is an "infinite dimensional" optimization problem
- Solutions are open-loop trajectories
- There are a handful of problems with analytic solutions but not many
- We will focus on the discrete-time setting in this class

* Discrete Time:

$$\min_{\substack{u \in U \\ x \in X \\ u_i \in U_i}} J(x_0, u_{1:N-1}) = \sum_{n=1}^{N-1} l(x_n, u_n) + l_f(x_N)$$

$$\text{s.t. } x_{n+1} = f(x_n, u_n)$$

$$u_{\min} \leq u \leq u_{\max} \leftarrow \text{"torque limits"}$$

$$c(x_n) \leq 0 \leftarrow \text{obstacle/safety constraints}$$

- This is a finite dimension

- Samples x_n, u_n are often called "knot points"
- Continuous \rightarrow discrete time using integration (e.g. Runge-Kutta)
- discrete \rightarrow continuous using interpolation

* Pontryagin's Minimum Principle

- Also called "maximum principle" if you maximize a reward
- First-order necessary conditions for a deterministic optimal control problem
- In discrete time, just KKT conditions

- Given :

$$\min_{\substack{x \in N \\ u \in U}} \sum_{n=1}^{N-1} l(x_n, u_n) + l_F(x_N)$$

$$\text{s.t. } x_{n+1} = f(x_n, u_n)$$

- We can form the Lagrangian

$$L = \sum_{n=1}^{N-1} l(x_n, u_n) + \lambda_n^T (f(x_n, u_n) - x_{n+1}) + l_F(x_N)$$

- This result is usually stated in terms of the "Hamiltonian"

$$H(x, u, \lambda) = l(x, u) + \lambda^T f(x, u)$$

- Plug it into L :

$$L = H(x_1, u_1, \lambda_1) + \left[\sum_{n=2}^{N-1} H(x_n, u_n, \lambda_{n+1}) - \lambda_n^T x_n \right] + l_F(x_N) - \lambda_N^T x_N$$

↑
Note change b indexing

- Take derivatives w.r.t. x and λ :

$$\frac{\partial L}{\partial \lambda_n} = \frac{\partial H}{\partial \lambda_n} - x_{n+1} = f(x_n, u_n) - x_{n+1} = 0$$

$$\frac{\partial L}{\partial x_n} = \frac{\partial H}{\partial x_n} - \lambda_n^T = \frac{\partial l}{\partial x_n} + \lambda_{n+1}^T \frac{\partial f}{\partial x_n} - \lambda_n^T = 0$$

$$\frac{\partial L}{\partial x_N} = \frac{\partial l_F}{\partial x_N} - \lambda_N^T = 0$$

- For u we write the min explicitly to handle torque limits:

$$u_n = \underset{u}{\operatorname{argmin}} H(x_n, u, \lambda_{n+1})$$

$$\text{s.t. } u \in \mathcal{U}$$

 Shorthand for "in feasible set"

e.g. $u_{\min} \leq u \leq u_{\max}$

- In Summary:

$$x_{n+1} = D_x H(x_n, u_n, \lambda_{n+1})$$

$$\lambda_n = D_\lambda H(x_n, u_n, \lambda_{n+1})$$

$$u_n = \underset{u}{\operatorname{argmin}} H(x_n, u, \lambda_{n+1})$$

s.t. $u \in \mathcal{U}$

$$\lambda_n = \frac{\partial l_F}{\partial x_n}$$

- These can be stated in continuous time:

$$\dot{x} = D_x H(x, u, \lambda)$$

$$\dot{\lambda} = D_\lambda H(x, u, \lambda)$$

$$u = \underset{u}{\operatorname{argmin}} H(x, u, \lambda)$$

s.t. $u \in \mathcal{U}$

$$\lambda(t_F) = \frac{\partial l_F}{\partial x}$$

* Some Notes:

- Historically many algorithms were based on integrating the continuous ODEs forward and backward to do gradient descent
- These are called "indirect" and/or "shooting" methods
- In continuous time, $\lambda(t)$ is called "costate" trajectory
- These methods have largely fallen out of favor as computers have improved.