

Last Time:

- LQR as a QP
- Riccati

Today:

- Infinite-Horizon LQR
 - Controllability
 - Dynamic Programming
-

* Infinite-Horizon LQR

- For time-invariant LQR, K matrices converge to constant values
- For stabilization problems we usually use constant K
- Backward recursion for P :

$$K_n = (R + B^T P_{n+1} B)^{-1} B^T P_{n+1} A$$

$$P_n = Q + A^T P_{n+1} (A - BK_n)$$

- Infinite-horizon limit $\Rightarrow P_{n+1} = P_n = P_{\text{inf}}$
 \Rightarrow solve as a root-finding / fixed-point problem
- Julia/MATLAB/Python dare function does this for you

\$\\$ Controllability

- How do we know if LQR will work?
- We already know $Q \geq 0$, $R > 0$
- For the time-invariant case, there is a simple answer
- For any initial state x_0 , x_N is given by:

$$\begin{aligned}x_N &= Ax_{N-1} + Bu_{N-1} \\&= A(Ax_{N-2} + Bu_{N-2}) + Bu_{N-1} \\&\quad \vdots \\&= A^Nx_0 + A^{N-1}Bu_0 + A^{N-2}Bu_1 + \dots + Bu_{N-1} \\&= \underbrace{[B \quad AB \quad A^2B \quad \dots \quad A^{N-1}B]}_C \begin{bmatrix} u_{N-1} \\ u_{N-2} \\ \vdots \\ u_0 \end{bmatrix} + A^Nx_0.\end{aligned}$$

- This is equivalent to a least-squares problem for $u_{0:N-1}$:

$$\begin{bmatrix} u_{N-1} \\ u_{N-2} \\ \vdots \\ u_0 \end{bmatrix} = \underbrace{[C^T(CC^T)^{-1}]_n}_{\text{"Pseudo-inverse"}} (x_N - A^Nx_0)$$

- For CC^T to be invertible:

$$\Rightarrow \text{rank}(CC^T) = n, \quad n = \dim(x)$$

- I can stop at n time steps in C because the Cayley-Hamilton theorem says that A^n can be written in terms of a linear combination of lower powers of A up to n

$$A^n = \sum_{n=0}^{n-1} \alpha_n A^n \quad (\text{for some } \alpha_n)$$

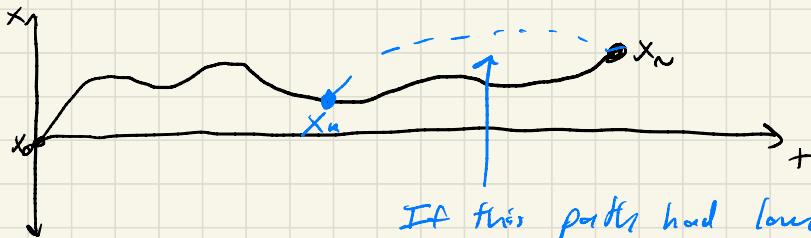
- Therefore adding more time steps/columns to C can't increase the rank:

$$\Rightarrow C = [B \ AB \ \dots \ A^{n-1}B]$$

"Controllability Matrix"

* Bellman's Principle

- Optimal control problems have an inherently sequential structure
- Past control inputs only affect future states, future inputs can't affect past states.
- Bellman's Principle ("Principle of Optimality") states the consequences of this for optimal trajectories.



If this path had lower cost starting at x_n , I would have taken it starting from x_0

- Sub-trajectories of optimal trajectories have to be optimal for appropriately defined sub problems

* Dynamic Programming

- Bellman's Principle suggests starting from the end of the trajectory and working backwards
- We've already seen hints of this from Riccati
- Define "optimal cost-to-go" aka "Value function"
 $V_k(x)$
- Encodes cost incurred starting from state x at time k if we act optimally

- For LQR

$$V_N(x) = \frac{1}{2} x^T Q_N x = \frac{1}{2} x^T P_N x$$

- back up one step and calculate $V_{N-1}(x)$:

$$V_{N-1} = \min_u \frac{1}{2} x_{N-1}^T Q x_{N-1} + \frac{1}{2} u^T R u + V_N(A_{N-1} x_{N-1} + B u_{N-1}) \quad *$$

$$\Rightarrow \min_u u^T R u + (A x_{N-1} + B u_{N-1})^T P_N (A x_{N-1} + B u_{N-1}) \quad) D_u = 0$$

$$\Rightarrow R u + B^T P_N (A x_{N-1} + B u) = 0$$

$$\Rightarrow u_{N-1} = - \underbrace{(R + B_{N-1}^T P_N B_{N-1})^{-1} B_{N-1}^T P_N A_{N-1} x_{N-1}}_{K_{N-1}}$$

- plug $u = -K x$ back into *

$$\Rightarrow V_{N-1}(x) = \frac{1}{2} x^T \underbrace{[Q + K^T R K + (A - B K)^T P_N (A - B K)]}_{P_{N-1}} x$$

$$\Rightarrow V_{N-1}(x) = \frac{1}{2} x^T P_{N-1} x$$

- Now we have a backward recursion for K and P that we iterate until $K = 0$

* Dynamic Programming Algorithms:

$$V_n(x) \leftarrow l_n(x)$$

$K \leftarrow N$

while $K > 1$

$$V_{K-1}(x) = \min_{u \in U} [l(x, u) + V_n(f(x, u))]$$

end $K \leftarrow K-1$

- If we know $V_n(x)$, the optimal policy is:

$$u_n(x) = \arg\min_{u \in U} [l(x, u) + V_n(f(x, u))]$$

- DP equations written equivalently in terms of "action-value" or "Q" function:

$$S_n(x, u) = l(x, u) + V_{n+1}(f(x, u))$$

- Usually denoted $Q(x, u)$ but we'll use $S(x, u)$

- Avoids need for explicit dynamics model

* The Curse

- DP is sufficient for global optimum
- Only tractable for simple problems (LQR, low-dimensional)
- $V(x)$ stays quadratic for LQR but becomes impossible to write analytically even for simple nonlinear problems.
- Even if we could, $\min S(x,u)$ will be non-convex and possibly hard to solve.
- Cost of DP blows up with state dimension due to cost of representing $V(x)$

* Why do we care?

- Approximate DP with a function approximator for $V(x)$ or $S(x,u)$ is very powerful.
- Forms basis of lots of modern RL
- DP generalizes to stochastic problems (just wrap everything in expectations). Pontryagin does not.

* Finally : What are the Lagrange Multipliers ?

- Recall Riccati derivation from QP:

$$\lambda_u = P_u x_u$$

- From DP:

$$V(x) = \frac{1}{2} x^T P x$$

$$\Rightarrow \lambda_u = \nabla_x V_u(x)$$

- Dynamics multipliers are cost-to-go gradients

- Carries over to nonlinear setting (not just LQR)